

Northstar AI Delivery Guide

Document owner: AI Enablement. Version: 1.0.

The team ships on a two-week release cadence. A release candidate needs passing unit and integration tests, at least 85 percent line coverage, a dependency scan with no unresolved critical findings, and approval from one engineer who did not author the change. The release manager owns the go or no-go decision.

Routine releases are frozen at 16:00 UTC on the Wednesday before deployment. A failed quality gate blocks release; a deadline is not an exception.

Product experiments

New AI assistance features begin with 10 percent of eligible users for 24 hours. The primary success metric is task completion rate, with a target lift of at least 5 percent relative to the control. Each arm needs at least 1,000 eligible tasks before a ship decision.

The error-rate guardrail allows no more than a 0.5 percentage-point increase. The latency guardrail allows no more than a 10 percent increase in p95 latency. If either guardrail fails, exposure returns to zero while the team investigates. Segment results are reported for new and returning users.

Cost and operations

The monthly hosted-model budget is 1,200 US dollars. The finance alert fires at 70 percent of the budget and the engineering escalation fires at 90 percent. Cost reports separate embedding input, generation input, and generation output.

Repeated, identical document chunks may reuse cached embeddings when the model, dimension, and normalized text hash all match. Generated answers are not cached because freshness and experiment attribution are more important than the small latency saving. Usage dashboards retain aggregate metrics for 90 days and raw sanitized trial records for 30 days.